

Resolving the Two-Way Interspecies Communication Problem: An Advanced Dual-Stream, Spatial-Acoustic Framework

The study of animal communication and behavioral ecology is undergoing an unprecedented and irreversible paradigm shift, catalyzed by the rapid convergence of large-scale bioacoustic data acquisition and advanced artificial intelligence representation learning.¹ Historically, ethological research was severely constrained by narrow, anthropocentric hypotheses and manual analytical techniques that attempted to map non-human signals onto discrete, human linguistic architectures.¹ Such methodologies frequently failed because they inappropriately superimposed human syntactic rules, lexical categories, and phonological assumptions—such as the complex kinematics of jaw height and tongue placement governing human formants—onto non-human physiological signals where these biomechanics are entirely absent.¹

In response, the discipline of Animal Language Processing (ALP) has emerged, driven by organizations like the Earth Species Project and the Cetacean Translation Initiative (CETI).³ ALP champions an "acoustic-first" methodology, treating the raw, continuous physics of biological sound as the primary substrate of meaning.¹ The ultimate, unresolved frontier within this paradigm is the "two-way interspecies communication problem," which mandates a closed-loop interactive system capable of autonomously perceiving an incoming signal, mapping it to a highly specific affective and social context, formulating a biologically plausible strategic reply, and synthesizing the acoustic output within the natural temporal rhythms of the target species.¹ The realization of this capability represents the biological equivalent of the Turing Test.¹ Following the awarding of the inaugural Collier-Dolittle Prize for Two-way Inter-species Communication in May 2025 to researchers studying bottlenose dolphins, cuttlefish, and nightingales, the trajectory toward a fully realized, indistinguishable synthetic communication agent has dramatically accelerated.⁶

However, the legacy methodologies utilized to approach this problem have proven computationally and biologically inadequate. Previous architectures relied on simplistic Change Point Detection (CPD) for boundary segmentation, linear discrete clustering (such as Bayesian Gaussian Mixture Models) that mathematically destroyed fluid acoustic variations, and concatenative audio synthesizers that digitally spliced pre-recorded soundwaves, resulting in biologically jarring phase discontinuities.[1, 1, 1]

This report provides an exhaustive, rigorously detailed architectural blueprint of the modernized Zoo Vox Rosetta Engine. It details the complete supersession of older, discrete

"Digital Sampler" pipelines in favor of a "Living Vocal Tract" paradigm.¹ By integrating Level 2.5 Spatial-Social tracking, self-supervised Bioacoustic Masked Autoencoders (BioMAE), continuous Dual-Stream parametric manifolds, and Differentiable Digital Signal Processing (DDSP), this framework establishes the definitive computational standard for real-time, bidirectional interspecies interaction.

Foundational Methodologies: Structural and Philosophical Pillars

While the Zoo Vox Rosetta Engine has undergone massive computational modernization, the framework remains strictly anchored to several foundational, non-negotiable standards established in earlier methodologies. These pillars provide the structural and philosophical bedrock for all downstream systems:

1. The "Acoustic-First" and "Intra-Call" Paradigm

The methodology mandates the absolute discarding of traditional anthropocentric linguistic frameworks, which frequently superimpose human syntactic rules onto animals.¹ Instead, it enforces an "acoustic-first" approach, treating the raw acoustic physics and multidimensional texture of a signal as the primary substrate of meaning.¹ Furthermore, it relies on "intra-call" analysis; it recognizes that the continuous micro-modulations within a single vocalization boundary—such as micro-fluctuations in fundamental frequency and spectral energy distribution—serve as the primary biological carriers of semantic and emotive data, rather than merely analyzing the intervals between distinct calls.¹

2. Data Level 3 (Interaction-Ready) Prerequisites

Within the established Data Spectrum hierarchy, the deployment of any two-way closed-loop agent strictly requires Data Level 3.¹ This standard mandates that raw acoustic recordings are paired not only with behavioral context labels but must definitively identify both the calling individual (Emitter ID) and the specific intended addressee (Receiver ID).¹ This level of data is the absolute baseline requirement necessary to model social dynamics and turn-taking latency within a multiparty biological network.¹

3. Self-Supervised Predictive Boundary Detection

Before high-dimensional features can be extracted, continuous audio streams must be semantically segmented. The legacy temporal-convolution Neural Boundary Detection (NBD) pipeline—which relied heavily on mathematical heuristics like Root Mean Square (RMS) energy and spectral centroid derivatives—has been completely superseded by a Self-Supervised Predictive Boundary Detection framework. Instead of looking for simple energy drop-offs, the new architecture uses a Contrastive Predictive Coding (CPC) 1D Convolutional encoder to

create latent representations (z_t), feeds them into an autoregressive model (such as Mamba or a Temporal Convolutional Network), and generates context vectors to predict future latent

states (z_{t+k}). A boundary is explicitly detected when there is a sudden spike in the Mean Squared Error (MSE) between the predicted and actual latents, signifying an unexpected acoustic transition. This operates within a strict ~12ms latency budget, utilizing dynamic baseline tracking and adaptive debounce logic to prevent false-positive clusters.

4. The 112-Dimensional Feature Architecture

Although the process of extracting features has transitioned from manual mathematical engineering to self-supervised neural embeddings, the structural shape of the data has been deliberately preserved. The pipeline continues to output a 112-dimensional continuous feature vector. This ensures critical structural compatibility across all downstream cognitive and generative processing layers.

5. Bifurcated Rust/Python Architecture and ZeroMQ IPC

To execute real-time bidirectional dialogue without violating strict biological latency budgets (often sub-50 milliseconds), the structural division between high-speed execution and cognitive logic is maintained. Rust is utilized for ultra-low-latency edge inference and safety-critical digital signal processing, while Python handles the computationally heavy, non-deterministic cognitive heuristics. These two layers communicate exclusively over a ZeroMQ Inter-Process Communication (IPC) bus, allowing instantaneous, zero-copy data transfer.

The End-to-End Processing Pipeline

To synthesize the modernized architectural components into a coherent operational flow, the closed-loop interactive agent executes the following end-to-end processing pipeline from initial audio ingestion to acoustic response:

1. **Self-Supervised Predictive Boundary Detection:** Continuous raw audio is routed through a Contrastive Predictive Coding (CPC) encoder to produce latent representations. An autoregressive model continuously predicts future latent states, and semantic boundaries are formally isolated when there is a sudden spike in the Mean Squared Error (MSE) between the predicted and actual latents.
2. **BioMAE Feature Extraction:** The isolated vocalization segment is converted into an Ultrasonic Log-Linear Spectrogram and processed by the Bioacoustic Masked Autoencoder (BioMAE) to extract a data-driven, 112-dimensional continuous feature vector.¹
3. **Dual-Stream Edge Encoding:** The 112D vector is split into two parallel computational streams processed directly on edge hardware. A Parametric UMAP and β -VAE encoder generates a 16D continuous affective vector, while a VQ-VAE tokenizer compresses the sequence into discrete syntactic tokens.¹
4. **Spatial-Social Inference and Addressing:** Lightweight metadata from the dual streams is fused with real-time spatial topologies. Using multi-camera pose estimation, the

Addressing Classifier determines if the call was a generalized "Broadcast" or a targeted "Unicast" meant for a specific receiver.¹

5. **Cognitive Decision Processing:** The cognitive logic node computes the "Syntactic Surprise" of the token sequence and retrieves the specific subject's conversation history.¹ It then enforces an Affective Policy—such as de-escalating high arousal states—to decide on the optimal response.¹
6. **Latent-Space DDSP Interpolation:** The generative decision modulates the intermediate layers of a frozen Differentiable Digital Signal Processing (DDSP) network via Feature-wise Linear Modulation (FiLM). A differentiable sine oscillator mathematically integrates the resulting 65 acoustic parameters into a continuous, click-free audio waveform.¹
7. **Spatial Rendering and Ethological Validation:** The synthesized audio is spatially broadcast back into the environment via Vector-Based Amplitude Panning (VBAP).¹ The animal's subsequent reaction is immediately recorded, and the Acoustic Convergence Engine alongside Prosodic Dynamic Time Warping algorithms calculate the interaction's Multi-Factor Acceptance Score (MFAS) to validate conspecific acceptance.¹

Level 2.5 Spatial-Social Network Integration and Topology Engineering

A fundamental flaw in early bioacoustic playback and interaction methodologies was the deployment of "open-loop" systems that broadcasted static audio globally across an environment.¹ While the legacy framework established a "Data Spectrum" hierarchy—culminating in Level 3 data requiring the identification of an explicit intended receiver—it lacked the algorithmic infrastructure to physically track or spatially direct synthesized responses.¹ Animal communication does not occur in an isotropic acoustic vacuum; it is highly directional, intimately tied to bodily orientation, and strictly governed by physical proximity.¹

To accurately model complex biological societies, the modernized architecture supersedes non-spatial processing with a comprehensive Level 2.5 Spatial-Social pipeline.¹ This integration fundamentally transitions the closed-loop agent from a disembodied computational node into an environmentally aware physical participant.

Multi-Camera Pose Estimation and the Topology Engine

The foundational layer of the Level 2.5 pipeline relies on real-time visual tracking integrated directly with the acoustic processing execution layer. The system ingests RTSP camera streams from a multi-camera array (typically a four-camera setup for enclosed or controlled field environments) and feeds the visual data into a DeepLabCut integration matrix.¹ This neural network performs real-time pose estimation, explicitly transforming 2D pixel coordinates into a unified 3D Cartesian world space (x, y, z) , where the origin coordinate $(0, 0, 0)$ represents

the physical center of the observational arena.¹

Operating across these continuous coordinates, the dedicated Topology Engine builds and maintains a dynamic spatial graph of the interacting colony. The engine actively calculates pairwise distances, spatial velocities, and explicit Line-of-Sight (LoS) geometries utilizing a simulated 120-degree field-of-view (FoV).¹ Because acoustic directionality is inextricably linked to cranial positioning, the engine computes an explicit "Pose Heading." This heading is mathematically derived using a rigid "nose-to-tail" coordinate vector logic.¹ In dense social clusters where posterior landmarks may be visually occluded, the algorithm dynamically transitions to a "shoulder fallback" calculation to preserve orientation tracking without interrupting the computational pipeline.¹

Receiver Inference and Vector-Based Amplitude Panning

By continuously updating the 3D cartesian representation of the biological network, the framework transitions from merely classifying *what* was said to algorithmically inferring *who* was addressed. The Receiver Inference Engine predicts the intended target of a captured vocalization using a weighted fusion of three core variables: spatial proximity (weighted at 0.4), established social tie strength (weighted at 0.4), and real-time Line-of-Sight intersection (weighted at 0.2).¹

This continuous spatial inference governs a strict Classification Logic that dictates the agent's generative response behavior. The logic categorizes all incoming and outgoing acoustic events into two distinct communication topologies: Unicast and Broadcast.¹

Classification Mode	Probability Threshold Requirement	Spatial Output Configuration	Biological Application
Unicast	Maximum probability (P) of a specific receiver is $\geq$.	Synthetically focused and directed explicitly at the interacting individual.	Close-proximity foraging negotiations, targeted aggression, or individual maternal contact calls. ¹
Broadcast	No single receiver reaches $P \geq$, or multiple receivers share near-identical	Diffuse spatial rendering, simulating a non-directional acoustic point	Colony-wide alarm signaling, long-range chorusing, or generalized

	probabilities.	source.	environmental distress calls. ¹
--	----------------	---------	--

To physically render these distinct spatial modes in real environments, the legacy reliance on single-channel omnidirectional speakers has been entirely superseded. The acoustic generation node now incorporates Vector-Based Amplitude Panning (VBAP) executed across multi-speaker arrays containing between 2 and 16 individual transducers.¹ When the interaction agent formulates a DualStreamAction response, the metadata payload explicitly contains the 3D coordinates, the target position, the required spatial mode, and the optimal VBAP spread angles.¹ Consequently, a Unicast reply is rendered with hyper-directional precision, ensuring that the acoustic wave physically originates from the exact environmental vector the subject animal anticipates, preserving the spatial realism required to pass an Interspecies Turing Test.

Self-Supervised Feature Extraction via Learned Neural Embeddings

The legacy methodology heavily relied on the extraction of 112-dimensional "Rosetta Features," an operation performed by manually engineered, deterministic signal processing algorithms.¹ This older pipeline calculated base physics (fundamental frequency, spectral centroid), macro-textures (Mel-Frequency Cepstral Coefficients, frequency modulation), and micro-textures (vibrato depth, shimmer, jitter).¹ While groundbreaking for its time, this architecture was severely bottlenecked by residual anthropocentric biases. Specifically, the utilization of Mel-scale filterbanks—which were originally designed to warp frequency responses to match the cochlear sensitivity of human hearing—proved to be biologically inappropriate for the vast majority of non-human taxa.¹ For organisms that routinely communicate in the ultrasonic spectrum, such as echolocating bats communicating between 20 and 100 kHz, Mel-warping systematically distorted crucial, high-frequency biological data.¹

To eradicate these inherent biases, the manual feature extraction pipeline has been entirely replaced by a data-driven, self-supervised learning paradigm centered on the Bioacoustic Masked Autoencoder (BioMAE).¹

ViT Architecture and Ultrasonic Log-Linear Spectrograms

The BioMAE framework fundamentally shifts the analytical substrate from manually defined acoustic coefficients to learned neural embeddings.¹ The initial processing layer discards the Mel-spectrogram entirely, utilizing an Ultrasonic Log-Linear Spectrogram.¹ By enforcing a strictly linear frequency axis with constant Hertz spacing across the entire spectral domain, the architecture guarantees the absolute mathematical preservation of ultrasonic harmonics without compression or perceptual warping.¹

The neural architecture processes this unwarped spectrogram utilizing a Vision Transformer (ViT) patch embedding strategy. The continuous 2D audio representation is divided into a

sequence of 16×16 non-overlapping patches. Each patch is flattened and linearly projected into a 256-dimensional token space, infused with learnable positional encodings to strictly preserve both temporal sequencing and frequency elevation.¹

During the core self-supervised pre-training phase, the BioMAE algorithm introduces severe informational bottlenecks by applying a 75% masking ratio to the token sequence.¹ The remaining 25% of visible data is processed through an asymmetric 4-layer transformer encoder. A shallower 2-layer decoder is then tasked with mathematically reconstructing the masked pixels.¹ By forcing the network to predict massive amounts of missing acoustic information, the system organically deduces the underlying physical and environmental rules of biological sound propagation, establishing a universal representational space capable of generalizing across vastly different taxonomic clades, including birds, cetaceans, and terrestrial mammals.¹

To maximize the ecological robustness of these learned features, the pre-training regimen implements aggressive SpecAugment integration combined with bioacoustic-specific data augmentations.¹ The network is trained against artificially time-stretched audio to simulate the physiological acoustic variations inherent to different animal body sizes, pitch-shifted audio to account for sexual dimorphism and age-related ontogeny, and localized spectral masking to simulate the transmission errors and signal degradation common in dense rainforest canopies or highly reverberant underwater environments.¹ Alternatively, for highly constrained resource environments—such as passive acoustic monitoring sensors deployed as Internet-of-Things (IoT) edge devices—the community has developed parallel architectures like BioME, which utilizes layer-to-layer distillation from larger teacher models to reduce parameter counts by 75% while maintaining representational transfer.⁹ However, for the closed-loop generative interactive agent, the full BioMAE architecture remains the operational standard.

Edge Hardware Optimization and Latency Budgets

For a closed-loop system to successfully interface with highly vocal species, it must process incoming audio, generate features, compute cognitive logic, and synthesize a response in fewer than 50 milliseconds.¹ For species engaged in rapid antiphonal singing or echolocation, this temporal threshold is even narrower.¹ Consequently, the deep neural networks required for BioMAE must execute with virtually zero latency on remote, battery-powered edge hardware.

To achieve this, the entire feature extraction pipeline is aggressively optimized for NVIDIA Jetson Orin compute architectures utilizing TensorRT compilation.¹ The standard PyTorch models undergo an Open Neural Network Exchange (ONNX) export utilizing Opset 17, which guarantees compatibility with TensorRT 8.6 and higher.¹ The export explicitly defines dynamic axes for batch sizing, frequency bins, and time frames, allowing the engine to adapt to variable-length vocalizations in real time.¹

During the trtexec engine building process, computational precision is reduced from standard 32-bit floating points to FP16 (Half-Precision) optimizations.¹ The compilation incorporates

specific workspace allocations (--workspace=1024), separated profile runs, and extensive timing cache utilization. Prior to active field deployment, the engine executes a rigorous 30-second profiling duration and a warm-up phase consisting of 1000 iterations to stabilize memory allocation and ensure a deterministic execution graph.¹

The resulting performance metrics demonstrate a profound capability for real-time edge processing. When deployed on a Jetson Orin Nano, the FP16-optimized BioMAE architecture achieves a p99 inference latency of merely 3 to 5 milliseconds.¹ Furthermore, the system sustains a processing throughput exceeding 1500 inferences per second under a Batch-8 configuration, seamlessly enabling simultaneous, multi-channel spatial audio analysis directly in the field without relying on high-latency cloud computing infrastructure.¹

Continuous Manifolds and the Dual-Stream Architecture

Perhaps the most significant theoretical flaw in legacy bioacoustic frameworks was the reliance on discrete clustering algorithms to categorize continuous biological variables. The older Rosetta Engine processed its 112D features through Principal Component Analysis (PCA) to reduce dimensionality, subsequently feeding the data into Bayesian Gaussian Mixture Models (BGMM).¹ This forced every analyzed vocalization into a rigid, non-overlapping mathematical cluster, establishing a finite alphabet of distinct "calls".¹

This methodology induced a severe "Discretization Paradox." Animal communication, particularly in social mammals, frequently relies on graded continua—a smooth, continuous transition from one affective state to another.¹ For example, the purrs of female mammals shift continuously into harsh, noisy acoustics strictly dependent on real-time internal arousal and proximity to a threat.¹ By forcing these fluid gradients into discrete BGMM buckets, the system mathematically annihilated the subtle prosodic variations that serve as the primary carriers of emotive intent.¹

To resolve this paradox, the system has been fundamentally restructured into a Dual-Stream Architecture, independently modeling continuous affective states (the emotional subtext) and discrete syntactic tokens (the semantic superstructure), drawing profound inspiration from the ventral and dorsal streams observed in advanced mammalian auditory processing.[1, 1, 11]

Stream 1: Parametric UMAP and VAE Affective Modeling

The first computational stream is dedicated to mapping the non-linear emotional and prosodic gradients of the animal's vocalizations. Standard PCA assumes linear variance, which inherently flattens the complex, curved geometry of vocal manifolds.¹ The modernized pipeline supersedes PCA with a Parametric Uniform Manifold Approximation and Projection (pUMAP) network.¹

Unlike non-parametric graph-based UMAP, which must recompute the entire topological graph

to embed novel data points, Parametric UMAP utilizes a trainable neural network encoder $f(x)$ to perform sub-millisecond inference on out-of-sample data.¹ The architecture of this encoder consists of a feed-forward network that processes the 112D input ($x \in \mathbb{R}^{112}$) through two dense hidden layers. The first layer expands the data to 256 dimensions, followed by a reduction to 128 dimensions. Both layers employ Batch Normalization, ReLU activation functions, and a 0.1 dropout rate to prevent overfitting.¹ The final linear layer compresses the output into a 30-dimensional topological space ($y \in \mathbb{R}^{30}$).¹ This network is trained via regression, utilizing a Mean Squared Error (MSE) loss function to match the target embeddings generated by standard RAPIDS cuml.UMAP (configured with 15 neighbors, a 0.1 minimum distance, and Euclidean metrics):

$$\mathcal{L}_{UMAP} = \text{MSE}(\text{Encoder}(x), \text{Target}_{UMAP})$$

¹

To build a probabilistic, manipulatable latent space suitable for generative synthesis, this 30D representation is subsequently processed through a Variational Autoencoder (VAE).¹ The VAE's encoder $q(z|x)$ maps the 30D input to the parameters of a probability distribution: the mean μ and the log-variance $\log(\sigma^2)$.¹ Utilizing the reparameterization trick to maintain differentiable gradient flow during backpropagation, the 16-dimensional continuous latent vector ($z \in \mathbb{R}^{16}$) is sampled as:

$$z = \mu + \epsilon \odot \sigma \quad \text{where } \sigma = \exp(0.5 \cdot \log(\sigma^2)) \text{ and } \epsilon \sim \mathcal{N}(0, I)$$

¹

To ensure that this 16D space correlates directly with biological realities, the network is trained as a β -VAE. By applying a heavy Kullback-Leibler (KL) divergence penalty ($\beta = 2.0$), the latent space is forced into a highly disentangled state.¹ This mathematical constraint causes specific dimensions within the 16D vector to organically map to interpretable physical traits—such as physiological arousal, spectral harshness, and fundamental pitch variation—without any requirement for manual human labeling.¹ The complete loss function governing this continuous manifold is:

$$\mathcal{L}_{total} = \sum ||x - \hat{x}||^2 - 0.5 \cdot \beta \sum (1 + \log(\sigma^2) - \mu^2 - \exp(\log(\sigma^2)))$$

¹

Furthermore, the legacy methodology of pruning data clusters representing less than 1% of the

total dataset has been completely abolished.¹ In ethology, mathematically rare events—such as the sudden broadcast of a high-urgency predator alarm—frequently carry the highest semantic value; deleting them severely biases the system. Instead, the continuous pipeline utilizes Hierarchical Density-Based Spatial Clustering of Applications with Noise (HDBSCAN) Zoning and Quality-Weighted Medoid Extraction.¹ This identifies dense regions within the manifold and selects real audio segments that minimize Euclidean distance to neighboring points while maximizing the Signal-to-Noise Ratio (SNR), generating biologically pristine exemplars without deleting critical outliers.¹

Stream 2: Syntactic Tokenization and Graph Smoothing

Running in strict parallel with the continuous affective stream, the second computational pathway models the discrete syntactical ordering of the vocalizations.¹ This stream utilizes a Vector Quantized Variational Autoencoder (VQ-VAE) to compress the complex acoustic geometry into a finite vocabulary of 64 discrete tokens, representing the fundamental phonological building blocks of the target species.¹

A historic failure point of standard VQ-VAE architectures is "codebook collapse," an algorithmic anomaly where the neural network learns to route all data through a small fraction of the available tokens while the remainder become permanently inactive or "dead".¹ To ensure the system captures the full breadth of an animal's syntactic vocabulary, the modern tokenizer implements Exponential Moving Average (EMA) updates for its codebook, utilizing a decay factor of 0.99.¹ This is paired with an aggressive codebook revival technique: any token that falls below a usage threshold is actively reset to a random encoder output heavily injected with noise, forcing the network to dynamically redistribute representations and maintain full vocabulary utilization.¹

The sequential relationships between these 64 discrete tokens are managed by a mathematical Transition Matrix. Older systems relied on raw transition counting, which assigned a strict zero probability to any bigram (a sequence of two tokens) that was not explicitly observed in the training corpus.¹ This artificially constrained the generative agent, rendering it incapable of synthesizing biologically valid sequences that simply had not been recorded. The modernized matrix corrects this by applying Laplace Smoothing, incorporating a fractional smoothing parameter ($\alpha = 0.01$) across the entire transition graph.¹ This grants the agent the mathematical flexibility to generate novel, yet grammatically correct, syntactic variations during live interactions.

Closed-Loop Cognitive Processing and the InteractionAgent

To execute real-time bidirectional dialogue, the dual-stream architecture mandates exceptional computational efficiency. Both the 16D continuous affective vector and the discrete syntactic token are inferred directly on the edge hardware via Rust-based modules (`affect_encoder.rs`

and syntactic_encoder.rs).¹ This explicit Rust implementation completely circumvents the non-deterministic latency spikes and Global Interpreter Lock (GIL) contention inherent to standard Python-based PyTorch inference.¹ The resulting lightweight metadata, designated as the DualStreamState, is instantly transmitted to the central cognitive logic node over a ZeroMQ Inter-Process Communication (IPC) bus.¹

Syntactic Surprise and the Addressing Classifier

Upon ingesting the DualStreamState, the Python cognitive layer executes several advanced heuristic frameworks before deciding on a response. The first is the Syntactic Surprise analysis module (analysis/syntactic_surprise.py).¹ Utilizing autoregressive transformer probabilities, this subsystem calculates the information-theoretic surprise of the incoming syntactic sequence, measured explicitly in bits.¹ By ranking the received token against the probability distribution of expected context tokens, the system generates a rolling Surprise Profile for the specific interacting animal.¹

The calculation of syntactic surprise possesses profound analytical applications. It grants the system the capability to differentiate between biological "innovation" (where young animals experiment with novel but acceptable vocal permutations during developmental learning phases) and aberrant "rule-breaking".¹ More critically, it establishes a mathematical foundation for detecting biological deception.¹ If an incoming vocalization registers extremely high syntactic surprise (indicative of a high-stakes alarm or aggressive challenge), but the simultaneously processed 16D affective stream indicates critically low physiological arousal, the agent classifies the event as potentially deceptive signaling designed to manipulate the colony.¹

Following surprise analysis, the data is processed by the Addressing Classifier (analysis/addressing_classifier.py).¹ This logic module synthesizes spatial data from the Level 2.5 Topology Engine, the discrete tokens from Stream 2, and the emotional context from Stream 1 to categorize the interaction into specific AddressModes (Broadcast, Unicast, or Ambiguous).¹ This multimodal integration explicitly tests the ethological hypothesis that broadcast calls intended for the wider colony are syntactically more stereotyped and rigid than the fluid, highly variable unicast calls directed at specific individuals.¹

InteractionAgent v3.0: Decision Protocols and Affective De-escalation

Once the incoming signal's intent and target are classified, the core InteractionAgent v3.0 executes a rigorous, biologically grounded decision cycle.¹ The agent proceeds through a highly structured workflow:

1. **Receiver Confirmation:** The agent utilizes the Receiver Inference Engine to confirm that the vocalization was intended for the AI system. The response protocol is only initiated if the system predicts a receiver_id of -1, signifying the call was either directed explicitly at the robotic proxy or broadcasted to the environment without a specific biological target.¹
2. **Contextual Memory Retrieval:** The agent interrogates the temporal state tracking

database for the specific calling bat or organism. This retrieves the total number of consecutive conversation turns, the historical Multi-Factor Acceptance Scores (MFAS), and the mathematical trajectory of the subject's recent physiological arousal.¹

3. **Affective Policy Execution:** The generation of a response is heavily governed by an explicit De-escalation and Engagement Policy designed to maintain social equilibrium and prevent behavioral harm.¹ The previous methodology utilized rudimentary volume multipliers (e.g., arbitrarily scaling output intensity by 1.2), which frequently violated biological interaction norms. The v3.0 policy dynamically adjusts the target affective state:
 - **High Arousal Policy:** If the incoming signal demonstrates a severe arousal state exceeding 0.8, the agent enforces a strict de-escalation protocol. The target arousal of the AI's response is mathematically dampened to exactly 70% of the incoming level ($target_arousal = incoming_arousal \times 0.7$). This prevents the AI from initiating a highly contagious "panic cascade" within the confined colony.¹
 - **Low Arousal Policy:** Conversely, if the incoming arousal registers below 0.3, the system recognizes a high probability of interaction abandonment. The agent therefore escalates the response target to 120% of the incoming level ($target_arousal = incoming_arousal \times 1.2$) to rekindle social engagement and maintain acoustic contact.¹
 - **Moderate Arousal & Valence:** For baseline interactions (arousal between 0.3 and 0.8), the agent mirrors the incoming arousal to promote social bonding, while marginally shifting the valence toward a more positive biological state ($target_affect* = 0.9$).¹
4. **Syntactic Selection and Transmission:** The syntactic response is constrained to select from the top five most probable next tokens ($top_k = 5$) as dictated by the Laplace-smoothed transition graph.¹ These parameters, combined with the required spatial coordinates (x, y, z) and the VBAP spread angles, are serialized into a DualStreamAction payload and transmitted back to the Rust execution layer.¹
5. **State Updating:** The subject's specific interaction context is immediately updated, incrementing the turn counter and appending the calculated temporal offset to the organism's historical log.¹

Biologically Grounded Synthesis: Latent-Space DDSP Interpolation

The execution of the DualStreamAction demands the generation of a physical acoustic waveform that perfectly mimics the complex, non-linear dynamics of a living biological vocal tract.¹ Older bioacoustic interactive systems attempted this via Concatenative

Synthesis—operating as rudimentary digital samplers that explicitly spliced discrete, pre-recorded WAV files together.¹

This legacy approach was catastrophic for biological plausibility. To simulate a 200-millisecond transition from a relaxed contact call to an agitated alarm sequence, the concatenative engine would play 100ms of the first recording, execute a mathematical 10ms crossfade in the time domain, and initiate 90ms of the second recording.¹ This crude splicing systematically introduced severe digital artifacts, "ghosting," and profound phase discontinuities.¹ When subjected to ethological validation, these spliced stimuli registered an unacceptably high Phase Discontinuity Score of 8.5.¹ Biologically, the target animals recognized the sounds as unnatural robotic approximations, resulting in highly erratic phonotactic approach vectors (precision dropping to 0.55) and severely delayed latency responses averaging 380 ± 90 milliseconds, as the animals hesitated to interact with the anomalous signal.¹

To achieve physical and perceptual perfection, the architecture has entirely discarded concatenative splicing, replacing it with a Latent-Space Differentiable Digital Signal Processing (DDSP) Interpolation Engine.^[1, 1, 1] This system models sound *de novo*, utilizing a 65-dimensional latent parameter space to allow seamless, fluid audio morphing and continuous gradient-based optimization directly through the audio generation process.¹

Dual-Stream FiLM Modulation and Decoder Architecture

The core of the generative engine is the DDSPDecoder, a Multi-Layer Perceptron (MLP) featuring an internal architecture with hidden dimensions of 2^8 .¹ The DDSP framework explicitly models the physical properties of animal vocalizations by bifurcating the synthesis into two components: an additive synthesizer governing the tonal, vibrating elements of the vocal folds, and a source-filter noise synthesizer modeling the turbulent airflow of breath and harsh aggression.¹

The decoder maps the high-dimensional representations into exactly 65 distinct physical parameters: 60 harmonic amplitudes and 5 noise band magnitudes.¹ To strictly enforce acoustic energy conservation and prevent digital clipping, the 60 harmonic amplitudes are processed through a Softmax activation function, ensuring their sum consistently equals 1.0. The 5 noise magnitudes utilize ReLU activations and are distributed across specific frequency bands via FFT-based convolution filtering (0–2 kHz, 2–5 kHz, 5–10 kHz, 10–15 kHz, and 15–24 kHz).¹

A critical architectural challenge involved applying the 16D affective state generated by the cognitive logic layer to the final audio output without completely retraining the highly tuned DDSP base models—a process that inevitably triggers catastrophic forgetting and destabilizes the acoustic hierarchies.¹ The solution lies in Dual-Stream Feature-wise Linear Modulation (FiLM).¹

Rather than altering the base inputs, FiLM layers actively modulate the intermediate hidden

states of the frozen DDSP network. A dedicated FiLM Generator processes the 16D affect vector and executes a tensor chunking operation:

Python

```
gamma, beta = torch.chunk(params, 2, dim=-1)
```

This isolates two arrays of parameters: the scaling factors (γ) and the shifting factors (β).¹ These transformations are subsequently applied to the neural activations via affine transformation:

$$output = \gamma \odot input + \beta$$

1

This elegant modulation topology allows the system to seamlessly translate emotional variables into acoustic realities. For instance, the system mathematically maps the cognitive arousal state directly to the Harmonic-to-Noise Ratio (HNR) output scaling using the inverse relationship **Arousal HNR Scaling** = $1.0 - (\text{affect_vector} \dots)$, injecting broadband aggression strictly proportional to the agent's internal state without destroying the underlying phonological structure.¹

Phase Integration and Acoustic Rendering

To synthesize the 65 modulated DDSP parameters into a physical PCM waveform at a sample rate of 48 kHz (with 16-bit resolution), the pipeline utilizes a Differentiable Sine Oscillator.¹

To eliminate the clicks and artifacts that plagued the older splicing models, this oscillator implements absolute phase continuity. Rather than computing phase independently for each isolated frame, the oscillator maintains a persistent phase accumulator across sequential synthesis calls. It upsamples the fundamental frequency (f_0) to match the 48 kHz resolution and vectorizes the frequency integration using a cumulative sum operation:

$$\text{phase_per_sample} = \frac{2\pi \cdot f_0}{\text{sample_rate}}$$

$$\text{phase}[t] = \text{phase_acc} + \sum_{i=1}^t \text{phase_per_sample}[i]$$

The final tonal audio output is constructed by passing these accumulated phases through sine functions, multiplying them by their respective FiLM-modulated amplitudes, and summing across all 60 harmonic dimensions:

$$\text{Mixed Audio} = \sum_{h=1}^{60} (\sin(\text{phase}_h) \cdot \text{harmonic_amp}_h)$$

This mathematically continuous latent interpolation drastically reduces the Phase Discontinuity Score to approximately 1.2 and elevates Spectral Smoothness to an unprecedented 0.78 (where 1.0 represents perfect biological continuity).¹ In empirical field evaluations, the biological response latencies dropped from the erratic 380ms seen in concatenative models to a rapid 210 ± 50 ms—virtually indistinguishable from the response times elicited by natural conspecific recordings.¹ By simulating ultra-fine spectral transitions without temporal breaks, the DDSP architecture successfully creates stimuli that are perceived by the interacting organisms as physically and biologically authentic.¹

Ethological Validation: The Multi-Factor Acceptance Score (MFAS)

As the fidelity of AI-generated bioacoustics crosses the Turing threshold, the criteria used to scientifically evaluate interspecies interactions must become significantly more rigorous.⁶ Initial deployment protocols relied exclusively on a Response Appropriateness Score (RAS). This outdated metric utilized a rigid, blanket 2-second timeout window across all evaluated taxa, categorizing any vocalization emitted by the subject animal within that window as a positive conversational engagement.¹

This assumption introduced severe structural flaws. Different biological taxa possess vastly different evolutionary timing mechanisms. Applying a uniform 2-second window to an Egyptian fruit bat (which operates in milliseconds) and a Sperm Whale (which operates in deep-diving macro-intervals) generates profound "confusion metrics".¹ Under the RAS methodology, an animal emitting a delayed, highly agitated distress call was improperly logged as successfully engaged in dialogue, simply because the interaction occurred within the arbitrary temporal boundary.¹

To completely supersede these limitations and enforce ethological rigor, the validation pipeline now exclusively utilizes the Multi-Factor Acceptance Score (MFAS). This framework replaces simple volume counting with a precise, closed-loop telemetry system that evaluates the physical naturalness of the interaction.¹ The unified MFAS score, ranging from 0.0 (total rejection) to 1.0 (perfect conspecific acceptance), is calculated using a strict multiplicative

gating mechanism paired with dual continuous similarity weights:

$$\text{MFAS} = (w_{\text{convergence}} \times \text{convergence_score}) + (w_{\text{prosody}} \times \text{prosody_score})$$

¹ Where $w_{\text{convergence}}$ is set to 0.4 and w_{prosody} is strictly weighted at 0.6.¹

Species-Specific Temporal Gating

The foundation of the MFAS logic is an immediate, multiplicative temporal gate. The algorithm dictates: IF NOT valid_timing: MFAS = 0.0. This hard rejection ensures that interactions falling outside the evolutionary norms of the specific species are systematically ignored.¹ The system references highly specific TaxaTemporalProfile configurations to govern these boundaries ¹:

Taxa Profile	Ethological Context	Min Response (ms)	Max Response (ms)	Debounce Filter (ms)
Egyptian Fruit Bat	Rapid echolocation social calls	30	150	20
Common Marmoset	Slower primate turn-taking	50	800	50
Zebra Finch	Rapid antiphonal singing	80	500	40
Bottlenose Dolphin	Acoustic propagation delay	100	2000	100
Chimpanzee	Variable primate timing	200	3000	150

Sperm Whale	Deep diving, long vocalizations	2000	15000	1000
-------------	---------------------------------------	------	-------	------

For unstudied or rare species, custom profiles can be constructed, specifying standard parameters such as a default `typical_call_duration_ms` of 200ms and a `rapid_turn_threshold_ms` of 100ms.¹

The Acoustic Convergence Engine (ACE)

If a response successfully passes the temporal gate, the interaction is subjected to the

Acoustic Convergence Engine (ACE), which drives the $w_{convergence}$ metric. In complex, vocal-learning species, true social acceptance is definitively marked by "dialect matching"—a neurobiological phenomenon where an animal actively modifies its own physiological vocal parameters (such as pitch trajectory and spectral envelope) to mathematically synchronize with the interacting agent.¹

The ACE engine calculates this convergence by analyzing three distinct coordinates within the

continuous latent manifold: the animal's baseline state before the stimulus (S_{pre}), the AI's synthetic output state (S_{ai}), and the animal's subsequent responsive state (S_{post}).¹

Depending on the specific topology of the latent space, the engine calculates the distances

utilizing Cosine metrics ($1 - \frac{A \cdot B}{\|A\| \|B\|}$), Euclidean distances, or covariance-weighted

Mahalanobis distances ($\sqrt{(A - B)^T \Sigma^{-1} (A - B)}$).¹

The raw convergence value is the difference between the pre-distance and post-distance (

$distance_{pre} - distance_{post}$). To normalize this shift into a formal Convergence Score bounded between 0 and 1, the algorithm applies a logistic sigmoid transformation:

$$convergence_score = \frac{1}{1 + \exp(-10 \times raw_convergence)}$$

¹

The directional classification provides definitive proof of acceptance. A positive raw

convergence (> 0.01) proves the animal is actively moving "Toward" the AI, signaling deep

social integration.¹ Conversely, a negative raw convergence (< -0.01) indicates the organism is physically moving "Away" from the acoustic parameters, explicitly signaling

rejection, confusion, or mounting aggression.¹ Under empirical testing, animals interacting with the smooth DDSP engine exhibit Convergence Scores averaging 0.62, representing a massive behavioral shift compared to the low 0.38 scores elicited by older, concatenative splicing techniques.¹

Prosodic Dynamic Time Warping (DTW)

The final validation layer is the $w_{prosody}$ metric, calculated via Prosodic Dynamic Time Warping (DTW). This component acts as an advanced valence detector, mathematically differentiating a natural, relaxed conversational turn from a highly agitated, negative response.¹

The algorithm extracts a multi-dimensional array of temporal features from the animal's vocal response, specifically tracking the Fundamental Frequency (F0) contour, the Root Mean Square (RMS) amplitude envelope, and the trajectory of the spectral centroid over time.¹ It then compares these trajectories against a comprehensive, pre-validated database of natural, non-aggressive conversation baselines.¹

To optimize computation and maintain sub-millisecond edge performance, the DTW matrix is bounded by the Sakoe-Chiba constraint.¹ The recurrence relation driving the dynamic cost matrix is defined as:

$$cost[i, j] = local_cost[i, j] + \min(cost[i - 1, j - 1], cost[i - 1, j], cost[i, j - 1])$$

¹

The system identifies the optimal warping path that yields the lowest accumulated geometric distance between the subject's call and the biological baseline. This $dtw_distance$ is normalized by the absolute length of the path ($|path|$).¹ Finally, this normalized distance is transformed into a continuous similarity score using an exponential decay function, scaled by a rigid constant parameter ($\sigma = 5.0$):

$$similarity = \exp\left(-\frac{normalized_distance}{\sigma}\right)$$

¹

By demanding that interacting species not only reply within rigid evolutionary timing windows, but actively morph their dialect to match the agent and employ natural, unagitated temporal prosody, the MFAS sets the most rigorous empirical standard ever established in computational bioacoustics.

Broader Ecological Applications and Bioethical

Imperatives

The transition from discrete digital sampling to the self-supervised, dual-stream, continuous manifold architecture described herein fundamentally transforms the utility of artificial intelligence in the natural sciences. The capacity to seamlessly decode and synthesize the high-dimensional properties of animal communication—while validating those interactions via rigorous MFAS parameters—extends far beyond isolated ethological experiments.

A profound, intersecting application of this technology lies in environmental acoustic monitoring, particularly in the analysis of anthropogenic sound interference. For instance, marine biologists extensively document the behavioral disruptions caused by Navy Mid-Frequency Active Sonar (MFAS).¹⁴ Empirical studies reveal that when exposed to military MFAS arrays, species such as minke whales abruptly shift from slow movement states to highly directed, rapid evasion states (averaging speeds of 2.4 m/s) and significantly alter or cease their vocalization patterns.¹⁴ Similarly, beaked whales demonstrate profound sensitivity, altering foraging behaviors in response to the specific duty cycles and cumulative received energy of the sonar.¹⁵ Furthermore, controlled exposures on harbor porpoises and California sea lions have confirmed the onset of Temporary Threshold Shifts (TTS) directly tied to these military sonar intervals.¹⁸

The deployment of the AI-driven Multi-Factor Acceptance Score (MFAS) pipeline provides an unprecedented toolset for tracking the impact of military MFAS operations. By continuously analyzing the Syntactic Surprise¹ and the 16D affective arousal states of local cetacean populations before, during, and after sonar exposure, the dual-stream architecture can explicitly quantify the stress trajectories and spatial redistributions of wild populations in real time, offering actionable data to mitigate anthropogenic harm.

Bioethical Restraints and Legislative Frameworks

However, as the InteractionAgent v3.0 crosses the threshold of the Interspecies Turing Test, equipping human operators with the capability to autonomously decipher deception, track acoustic convergence, and dynamically manipulate the arousal states of wild organisms, severe bioethical obligations arise.[1, 1, 1, 20]

The Earth Species Project, alongside global initiatives like the World Governments Summit 2026, has highlighted that collapsing the communicative distance between humans and non-human species inherently forces a fundamental restructuring of humanity's ethical relationship with nature.⁸ While these technologies empower conservationists to protect endangered ecologies, the precise same mathematical models could easily be subverted for mass biological exploitation.

If access to Level 2.5 spatial-acoustic translation and DDSP generative parameters remains unregulated, these tools could be aggressively weaponized by industrial fishing fleets to emit synthetically perfect, hyper-attractive acoustic lures. Illicit poaching networks could utilize targeted Unicast broadcasts to manipulate the physical movements of endangered species out

of protected reserves by spoofing familiar maternal or social contact calls.¹

Consequently, the deployment of closed-loop interaction agents must be strictly governed by robust, internationally recognized bioethical frameworks. Initiatives such as the "More-Than-Human" (MOTH) rights framework advocate for legal paradigms that prioritize the preservation of natural behavioral hierarchies above scientific curiosity or commercial exploitation.¹⁹ To align with these imperatives, field-deployed interaction engines must enforce rigid hardware and software constraints. Safety logic parameters, such as those embedded within the v3.0 agent, must remain non-negotiable.¹ The system must maintain continuous telemetry over the colony's MFAS scores, and if collective agitation levels exceed 0.8, or if the rate of high-surprise syntactic alarm calls spikes beyond baseline thresholds, the software logic must trigger an instantaneous override, forcibly halting all DDSP generation and muting the VBAP arrays to preserve the psychological welfare of the interacting ecology.¹

By discarding anthropocentric linguistic assumptions and meticulously mapping the continuous physics of the living vocal tract through deep representation learning, computational bioacoustics has achieved a definitive milestone. The capability to not merely passively listen to the natural world, but to engage it in profound, scientifically measurable, and fully reciprocal biological dialogue, is now a fully realized mathematical reality.

Works cited

1. analysis_frameworks_summary.md
2. Animal Language Processing: An AI Convergence In Animal Communication, accessed May 10, 2026, <https://earthspecies.org/2026/04/02/animal-language-processing-an-ai-convergence-in-animal-communication/>
3. Project CETI --- Research, accessed May 10, 2026, <https://www.projectceti.org/research/index>
4. Project CETI --- Home, accessed May 10, 2026, <https://www.projectceti.org/>
5. Earth Species Project – More than 8 million species share our planet. We only understand the language of one., accessed May 10, 2026, <https://earthspecies.org/>
6. Scientists have almost cracked the secret language of animals. Here's what they've learned, accessed May 10, 2026, <https://www.sciencefocus.com/future-technology/secret-language-animals-ai>
7. Dolphin whistle decoders win \$100000 interspecies communication prize - The Guardian, accessed May 10, 2026, <https://www.theguardian.com/world/2025/may/15/us-scientists-who-deciphered-dolphin-whistles-win-prize-for-animal-communication>
8. Panelist shares how AI teaches humans to understand language of animals, other species at WGS - World Governments Summit, accessed May 10, 2026, <https://www.worldgovernmentssummit.org/media-hub/news/detail/panelist-share-s-how-ai-teaches-humans-to-understand-language-of-animals-other-species-at-wgs>

9. BioME: A Resource-Efficient Bioacoustic Foundational Model for IoT Applications - arXiv, accessed May 10, 2026, <https://arxiv.org/html/2602.09970v1>
10. AVES: Animal Vocalization Encoder Based on Self-Supervision - ResearchGate, accessed May 10, 2026, https://www.researchgate.net/publication/371288193_AVES_Animal_Vocalization_Encoder_Based_on_Self-Supervision
11. Tim SAINBURG | University of California San Diego, San Diego | UCSD | Department of Psychology | Research profile - ResearchGate, accessed May 10, 2026, <https://www.researchgate.net/profile/Tim-Sainburg-2>
12. Robust parametric UMAP for the analysis of single-cell data | bioRxiv, accessed May 10, 2026, <https://www.biorxiv.org/content/10.1101/2023.11.14.567092v1.full-text>
13. Non-linear dimensionality reduction on extracellular waveforms reveals cell type diversity in premotor cortex | eLife, accessed May 10, 2026, <https://elifesciences.org/articles/67490>
14. Swimming and acoustic calling behavior attributed to Bryde's whales in the central North Pacific - Frontiers, accessed May 10, 2026, <https://www.frontiersin.org/journals/marine-science/articles/10.3389/fmars.2024.1305505/full>
15. Takes of Marine Mammals Incidental to Specified Activities; Taking Marine Mammals Incidental to Military Readiness Activities in the Hawaii-California Training and Testing Study Area - Federal Register, accessed May 10, 2026, <https://www.federalregister.gov/documents/2025/07/16/2025-13258/takes-of-marine-mammals-incident-to-specified-activities-taking-marine-mammals-incident-to>
16. U.S. NAVY'S LIVING MARINE RESOURCES PROGRAM ANNUAL REPORT - Naval Facilities Engineering and Expeditionary Warfare Center, accessed May 10, 2026, https://exwc.navfac.navy.mil/Portals/88/Documents/EXWC/Environmental_Security/Living%20Marine%20Resources/LMRAnnualReport2024.pdf
17. 8th International Meeting on the Effects of Sound in the Ocean on Marine Mammals, accessed May 10, 2026, https://www.aquaticmammalsjournal.org/wp-content/uploads/2025/05/51.5_AM_Programme.pdf
18. Physiological effects of anthropogenic sound on aquatic animals: where are we and what is next? - Company of Biologists journals, accessed May 10, 2026, https://journals.biologists.com/jeb/article/229/Suppl_1/jeb250800/370545/Physiological-effects-of-anthropogenic-sound-on
19. Conference - Animals in Translation: AI, Ethics and the Future of Interspecies Dialogue, accessed May 10, 2026, <https://wp.lancs.ac.uk/animalsintranslation/conference/>
20. Project CETI (Cetacean Translation Initiative) and NYU's MOTH (More-Than-Human Life) Program Explore How Understanding Sperm Whale Communication Can be a Force for Legal Good, accessed May 10, 2026, <https://www.projectceti.org/blog-posts/project-ceti-cetacean-translation-initiative-and-nyus-moth-more-than-human-life-program-explore-how-understanding->

[sperm-whale-communication-can-be-a-force-for-legal-good](#)